

# MnistdigitCNN

November 14, 2023

```
[10]: import tensorflow as tf #import the required libraries
      from tensorflow.keras.datasets import mnist
      from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
```

```
[3]: (X_train, y_train), (X_test, y_test) = mnist.load_data() #Load the MNIST
      ↪dataset:
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>  
11490434/11490434 [=====] - 2s 0us/step

```
[4]: X_train = X_train.reshape(X_train.shape[0], 28, 28, 1)#Preprocess the data
      X_train = X_train.astype('float32') / 255
      X_test = X_test.astype('float32') / 255
```

```
[5]: y_train = tf.keras.utils.to_categorical(y_train, 10)#Convert the target
      ↪variable to one-hot encoded format:
      y_test = tf.keras.utils.to_categorical(y_test, 10)
```

```
[6]: model = Sequential()#Define the architecture
      model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
      model.add(MaxPooling2D((2, 2)))
      model.add(Conv2D(64, (3, 3), activation='relu'))
      model.add(MaxPooling2D((2, 2)))
      model.add(Conv2D(64, (3, 3), activation='relu'))
      model.add(Flatten())
      model.add(Dense(64, activation='relu'))
      model.add(Dense(10, activation='softmax'))
```

Metal device set to: Apple M1

```
2023-11-14 20:27:01.128053: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305]
Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel
may not have been built with NUMA support.
2023-11-14 20:27:01.128621: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271]
Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0
```

MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id: <undefined>)

```
[7]: model.compile(optimizer='adam', loss='categorical_crossentropy',  
    ↪metrics=['accuracy'])#Compile the model
```

```
[8]: model.fit(X_train, y_train, epochs=5, batch_size=64, validation_data=(X_test,  
    ↪y_test))#Train the model
```

Epoch 1/5

2023-11-14 20:27:45.500945: W

tensorflow/core/platform/profile\_utils/cpu\_utils.cc:128] Failed to get CPU frequency: 0 Hz

2023-11-14 20:27:45.800232: I

tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113] Plugin optimizer for device\_type GPU is enabled.

938/938 [=====] - ETA: 0s - loss: 0.1910 - accuracy: 0.9424

2023-11-14 20:27:59.046753: I

tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113] Plugin optimizer for device\_type GPU is enabled.

938/938 [=====] - 15s 13ms/step - loss: 0.1910 - accuracy: 0.9424 - val\_loss: 0.0538 - val\_accuracy: 0.9822

Epoch 2/5

938/938 [=====] - 12s 13ms/step - loss: 0.0534 - accuracy: 0.9833 - val\_loss: 0.0465 - val\_accuracy: 0.9869

Epoch 3/5

938/938 [=====] - 12s 12ms/step - loss: 0.0386 - accuracy: 0.9879 - val\_loss: 0.0302 - val\_accuracy: 0.9894

Epoch 4/5

938/938 [=====] - 12s 12ms/step - loss: 0.0285 - accuracy: 0.9914 - val\_loss: 0.0319 - val\_accuracy: 0.9886

Epoch 5/5

938/938 [=====] - 12s 13ms/step - loss: 0.0227 - accuracy: 0.9928 - val\_loss: 0.0399 - val\_accuracy: 0.9880

```
[8]: <keras.callbacks.History at 0x28e058970>
```

```
[9]: test_loss, test_acc = model.evaluate(X_test, y_test)#Evaluate the model:  
    print('Test accuracy:', test_acc)
```

313/313 [=====] - 2s 6ms/step - loss: 0.0399 - accuracy: 0.9880

Test accuracy: 0.9880000352859497


```
[11]: import matplotlib.pyplot as plt #Import the required library
```

```
[12]: history = model.fit(X_train, y_train, epochs=5, batch_size=64,
    ↪validation_data=(X_test, y_test))#Train the model while keeping track of the
    ↪training accuracy:
```

```
Epoch 1/5
938/938 [=====] - 12s 13ms/step - loss: 0.0197 -
accuracy: 0.9935 - val_loss: 0.0271 - val_accuracy: 0.9911
Epoch 2/5
938/938 [=====] - 12s 13ms/step - loss: 0.0168 -
accuracy: 0.9945 - val_loss: 0.0362 - val_accuracy: 0.9892
Epoch 3/5
938/938 [=====] - 12s 13ms/step - loss: 0.0138 -
accuracy: 0.9955 - val_loss: 0.0247 - val_accuracy: 0.9927
Epoch 4/5
938/938 [=====] - 12s 13ms/step - loss: 0.0116 -
accuracy: 0.9959 - val_loss: 0.0327 - val_accuracy: 0.9910
Epoch 5/5
938/938 [=====] - 12s 13ms/step - loss: 0.0106 -
accuracy: 0.9966 - val_loss: 0.0359 - val_accuracy: 0.9900
```

```
[13]: train_accuracy = history.history['accuracy']#Access the training accuracy
    ↪values from the history object
```

```
[14]: plt.plot(train_accuracy)
plt.title('Training Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.show()
```



[ ]: